

# Contents

<b>1</b>	<b>Basic Test Results</b>	<b>2</b>
<b>2</b>	<b>README</b>	<b>5</b>
<b>3</b>	<b>ex3.pdf</b>	<b>6</b>
<b>4</b>	<b>q1.sql</b>	<b>10</b>
<b>5</b>	<b>q2.sql</b>	<b>11</b>
<b>6</b>	<b>q3.sql</b>	<b>12</b>
<b>7</b>	<b>q4.sql</b>	<b>13</b>
<b>8</b>	<b>q5.sql</b>	<b>14</b>
<b>9</b>	<b>q6.sql</b>	<b>15</b>

# 1 Basic Test Results

```
1  Extracting Archive:
2  Archive:  /tmp/bodek.Wpquy9/db/ex3/mohammadgh/presubmission/submission
3      inflating: ex3.pdf
4      inflating: q1.sql
5      inflating: q2.sql
6      inflating: q3.sql
7      inflating: q4.sql
8      inflating: q5.sql
9      inflating: q6.sql
10     inflating: README
11
12 *****
13 ** Testing that all necessary files were submitted:
14 README:
15     SUBMITTED
16 ex3.pdf:
17     SUBMITTED
18 q1.sql:
19     SUBMITTED
20 q2.sql:
21     SUBMITTED
22 q3.sql:
23     SUBMITTED
24 q4.sql:
25     SUBMITTED
26 q5.sql:
27     SUBMITTED
28 q6.sql:
29     SUBMITTED
30
31 *****
32 ** Checking for correct README format:
33 Output:
34 CREATE TABLE
35 CREATE TABLE
36 CREATE TABLE
37
38 Inserting conferences.csv
39 Output:
40 COPY 75
41
42 Inserting institutions.csv
43 Output:
44 COPY 401
45
46 Inserting authors.csv
47 Output:
48 COPY 3810
49
50 Note: The output is capped at 100 characters.
51
52 Running q1.sql
53 Output:
54   region      | countrycount
55   -----+-----
56   africa      |             2
57   asia        |            18
58   australasia |             2
59   canada      |             1
```

```

60     europe      |          1
61     europe      |         27
62     southamerica |          4
63 (7 rows)
64
65
66
67 Running q2.sql
68 Output:
69     region      |          insavg
70 -----+-----
71     africa      | 1.5000000000000000
72     asia        | 6.0000000000000000
73     australasia | 9.5000000000000000
74     canada      | 32.0000000000000000
75     europe      | 1.0000000000000000
76     europe      | 8.0740740740740741
77     southamerica | 5.0000000000000000
78 (7 rows)
79
80
81
82 Running q3.sql
83 Output:
84     name
85 -----
86 Abhradeep Thakurta
87 Aditya Akella
88 Adrian Perrig
89 Ahmed Eldawy
90 Alastair F. Donaldson
91 Alejandro Russo
92 Alessandro Orso
93 Alexander Aiken
94 Alex C. Snoeren
95 Ali Mesbah 0001
96 Anand Sivasubramaniam
97 Andreas Moshovos
98 Andreas Podelski
99 Andrew C. Myers
100 Anthony K. H. Tung
101 Antonio Filieri
102 Antony L. Hosking
103 Benjamin C. Pierce
104 Bingsheng He
105 Bin Ren
106 Björn Franke
107 Boris Glavic
108 Carlos Maltzahn
109 Chris Parnin
110 Christian S. Jensen
111 Christopher W.
112
113 Running q4.sql
114 Output:
115     region      | country | totalcount
116 -----+-----+-----
117     africa      | za      | 2.0
118     asia        | cn      | 498.0
119     australasia | au      | 121.0
120     canada      | ca      | 219.0
121     europe      | it      | 4.0
122     europe      | de      | 324.0
123     southamerica | br      | 18.0
124 (7 rows)
125
126
127

```

```

128 Running q5.sql
129 Output:
130         name
131 -----
132 Aapo Hyvärinen
133 Aaron Roth 0001
134 Aaron Schulman
135 Aaron Sidford
136 Aarti Gupta
137 Aarti Singh
138 Abdallah Saffidine
139 Abhik Roychoudhury
140 Abhinav Gupta 0001
141 Abhi Shelat
142 Abhradeep Thakurta
143 Abraham Bernstein
144 Abusayeed Saifullah
145 Achi Brandt
146 Achim Jung
147 Adam D. Smith
148 Adam Herout
149 Adam Klivans
150 Aditya Akella
151 Aditya G. Parameswaran
152 Aditya Grover
153 Aditya Kanade
154 Aditya P. Mathur
155 Adriana Iamnitchi
156 Adrián Jarabo
157 Adrian Perrig
158
159
160 Running q6.sql
161 Output:
162         name
163 -----
164 Aaron Sidford
165 Adam D. Smith
166 Adrian Vetta
167 Alexander A. Sherstov
168 Alexander Rakhlin
169 Angelika Steger
170 Benny Chor
171 Bernard Chazelle
172 Boaz Barak
173 Byung-Gon Chun
174 C. Greg Plaxton
175 Chang Xu 0002
176 Charles X. Ling
177 Chien-Ju Ho
178 Daniel C. Alexander
179 David Steurer
180 David Zuckerman
181 Derek Nowrouzezahrai
182 Dimitris Fotakis
183 Eli Upfal
184 Glencora Borradaile
185 Guillermo Sapiro
186 Hao Chen 0003
187 Hong Cheng 0001
188 Huy L. Nguyen
189 Jack Snoeyink
190 Jie Shao
191

```

## 2 README

1 mohammadgh,muaz.abdeen

## (67506) Databases – Spring 2022 – Exercise (3)

Muaz Abdeen 300575297

Mohammad Ghanayem 208653220

### Question (2): Indexing

*A* := authors (name, conference, year, institution, count, adjustedcount)

*C* := conferences (conference, area, subarea)

*I* := institutions (institution, region, country)

1.

(a)

```
SELECT DISTINCT year
FROM authors
WHERE institution='Hebrew University of Jerusalem';
```

(b)

#### QUERY PLAN

```
-----
HashAggregate  (cost=3904.75..3905.28 rows=53 width=4)
              (actual time=19.875..19.905 rows=43 loops=1)
    Group Key: year
    -> Seq Scan on authors  (cost=0.00..3902.30 rows=978 width=4)
                          (actual time=0.428..19.060 rows=929 loops=1)
    Filter: ((institution)::text = 'Hebrew University of Jerusalem'::text)
    Rows Removed by Filter: 163895
    Planning Time: 0.100 ms
    Execution Time: 19.964 ms
(7 rows)
```

Here the planner has decided to use a two-step plan: the child plan node scans the entire table sequentially, this is the meaning of Seq Scan, and then the upper plan node HashAggregate groups the records into a temporary hash table using the attribute year as a Group Key.

The grouping cost estimate (cost=3904.75..3905.28 rows=53 width=4) means that Postgres expects that the grouping will start at cost 3904.75 and finishes at 3905.28, the difference is the estimated cost in an arbitrary unit of computation to perform this operation. rows is the estimated number of rows this HashAggregate will return, and width is the estimated size in bytes of the returned

rows. The actual estimate (actual time=0.071..0.074 rows=4 loops=1) tell us actual time involved in execution. Here we have a new value loop which says the entire table was scanned one time.

The Seq Scan estimated cost (cost=0.00..3902.30 rows=978 width=4) is 3902.30 units of computation, starting at 0.00. where the actual cost was (actual time=0.428..19.060 rows=929 loops=1). The WHERE clause has been applied as a “Filter” condition attached to the Seq Scan plan node. This means that the plan node checks the condition for each row it scans, and outputs only the ones that pass the condition.

(c) CREATE INDEX ON authors(institution);

(d)

```
QUERY PLAN
-----
HashAggregate (cost=1636.41..1636.94 rows=53 width=4)
  (actual time=1.694..1.728 rows=43 loops=1)
    Group Key: year
    -> Bitmap Heap Scan on authors (cost=32.00..1633.96 rows=978 width=4)
          (actual time=0.107..0.882 rows=929 loops=1)
            Recheck Cond: ((institution)::text = 'Hebrew University of
                          Jerusalem'::text)
            Heap Blocks: exact=53
            -> Bitmap Index Scan on authors_institution_idx
                  (cost=0.00..31.76 rows=978 width=0)
                  (actual time=0.093..0.093 rows=929 loops=1)
                  Index Cond: ((institution)::text = 'Hebrew University of
                          Jerusalem'::text)
Planning Time: 0.124 ms
Execution Time: 1.799 ms
```

Here the planner has decided to use a three-step plan: the child plan node visits an index to find the locations of rows matching the index condition, and then the intermediate plan node actually fetches those rows from the table itself. Fetching rows separately is much more expensive than reading them sequentially, but because not all the pages of the table have to be visited, this is still cheaper than a sequential scan. The most upper plan is the grouping one as before.

Notice that the actual time using our defined index is 1.799 ms, which is almost tenth of the cost of the query without this index 19.964 ms.

2.a. (1) Without index we have to scan the entire table sequentially, that is to read all the blocks where the table is stored. Each block contains  $\left\lfloor \frac{2000}{180} \right\rfloor = 11$  rows, then the `authors` table takes  $\left\lceil \frac{12000}{11} \right\rceil = 1091$  blocks to store.

So, the query cost is: **1091**.

(2) Optimal branching factor is:  $d = \left\lfloor \frac{b+s}{p+s} \right\rfloor = \left\lfloor \frac{2000+8}{8+8} \right\rfloor = 125$ .

(3) Step 1: Find first relevant leaf:  $\left\lceil \log_{\left\lfloor \frac{d}{2} \right\rfloor} N \right\rceil = \left\lceil \log_{63} 12,000 \right\rceil = 3$ .

Step 2: Read all relevant leaves: since count values uniformly distributed over  $[1,20]$  range then there are:  $\frac{12,000}{20} = 600$  rows corresponding to `count=2`, then there are  $\left\lceil \frac{m}{\left\lfloor \frac{d}{2} \right\rfloor - 1} \right\rceil = \left\lceil \frac{600}{62} \right\rceil = 10$  leaves matching these rows.

Total cost = **13**.

2.b. (1) As the previous section, we have to traverse the entire tree, read all blocks, the query cost is: **1091**.

(2) The same as the previous section,  $d = 125$ .

(3) In addition to the step 1 and step 2 in the previous section, there is a step 3 here: to access the matching rows, we have to read all the **600** blocks containing the matching rows. The total = **613**.

2.c. (1) Scan the entire table: **1091**.

(2) Optimal branching factor is:  $d = \left\lfloor \frac{b+s}{p+s} \right\rfloor = \left\lfloor \frac{2000+26}{8+26} \right\rfloor = 59$ .



- (3) First, we traverse the tree down to the first matching leaf:  $\lceil \log_{[30]} 12,000 \rceil = 3$  . Second, we traverse all matching leaves, notice the table have at most  $\left\lceil \frac{12,000}{20 \times 80} \right\rceil = 7$  matching rows (maybe less, depends on how many year=1999 values are there), which fit in  $\left\lceil \frac{7}{29} \right\rceil = 1$  one leaf.

Total cost = **4** .

2.d. (1) Optimal branching factor is:  $d = \left\lfloor \frac{b+s}{p+s} \right\rfloor = \left\lfloor \frac{2000+22}{8+22} \right\rfloor = \mathbf{67}$  .

- (2) We have WHERE with a disjoint condition “or”, so we have to scan the tree two times, one for each value, the calculations for both are the same.

Traversing the tree down costs:  $\lceil \log_{[34]} 12,000 \rceil = 3$  , traversing the matching leaves costs:  $\left\lceil \frac{150}{33} \right\rceil = 5$  , where  $150 = \left\lceil \frac{12,000}{80} \right\rceil$  since the conference values uniformly distributed over 80 values. Now, since we are also indexing on name, there is no need to read the block containing the row.

The total cost for both conditions is  $= 2 \times (3 + 5) = \mathbf{16}$  .

2.e. (1) Optimal branching factor is:  $d = \left\lfloor \frac{b+s}{p+s} \right\rfloor = \left\lfloor \frac{2000+8}{8+8} \right\rfloor = \mathbf{125}$  .

- (2) Tree traverse until leaf costs at most  $\lceil \log_{[63]} 12,000 \rceil = 3$  .

There are 19 count values to scan, which spans  $\frac{12,000}{20} \times 19 = 11,400$  rows, so the matching rows will be in at most  $\left\lceil \frac{11400}{62} \right\rceil = \mathbf{184}$  leaves.

There are approximately 11,400 matching rows, each may be in a different block, but we will go over each block of the table at most once = **1091** .

Total  $= 3 + 184 + 1091 = \mathbf{1278}$  .

## 4 q1.sql

```
1 SELECT region, COUNT(DISTINCT country) AS countryCount
2 FROM institutions
3 GROUP BY region
4 ORDER BY region;
```

## 5 q2.sql

```
1  SELECT region, COUNT(DISTINCT institution) * 1.0 / COUNT(DISTINCT country) AS insAvg
2  FROM institutions
3  GROUP BY region
4  ORDER BY region;
```

## 6 q3.sql

```
1  SELECT DISTINCT name
2  FROM authors
3  WHERE conference in (SELECT conference
4                        FROM conferences
5                        WHERE area='systems')
6  GROUP BY name
7  HAVING SUM(count) > 1
8  EXCEPT
9  SELECT DISTINCT name
10 FROM authors
11 WHERE conference in (SELECT conference
12                      FROM conferences
13                      WHERE area='systems')
14 GROUP BY name
15 HAVING MAX(year) < 2014
16 ORDER BY name;
```

## 7 q4.sql

```
1  WITH PaperCounts(region, country, totalCount) as
2      (SELECT region, country, SUM(count) + 0.0
3       FROM institutions NATURAL JOIN authors
4       GROUP BY region, country)
5  SELECT P1.region, P1.country, P1.totalCount
6  FROM PaperCounts P1
7  WHERE P1.totalCount >= all (SELECT P2.totalCount
8                             FROM PaperCounts P2
9                             WHERE P1.region = P2.region)
10 ORDER BY region, country;
```

## 8 q5.sql

```
1  WITH NotSpecConfs(conference) AS
2      (SELECT conference
3       FROM authors
4       GROUP BY conference
5       HAVING COUNT(DISTINCT year) < 10)
6  SELECT DISTINCT name
7  FROM authors
8  EXCEPT
9  SELECT DISTINCT name
10 FROM authors
11 WHERE conference IN (SELECT * FROM NotSpecConfs)
12 ORDER BY name;
```

## 9 q6.sql

```
1  WITH RECURSIVE T(name, dist) AS (  
2      VALUES('Noam Nisan', 0)  
3      UNION  
4      SELECT second, dist + 1  
5      FROM T INNER JOIN (SELECT A1.name AS first, A2.name AS second  
6                          FROM authors A1 INNER JOIN authors A2  
7                          ON (A1.name != A2.name and  
8                              A1.conference = A2.conference and  
9                              A1.year = A2.year)) D  
10         ON (T.name = D.first)  
11         WHERE dist < 2  
12 )  
13 SELECT DISTINCT name FROM T ORDER BY name;
```